

Heterogeneous Parallelism Mini Project

Title: Implementation of High Performance, Lock Free and Concurrent Data Structures

Team 1

Madhav Jivrajani (PES1201800028)

M S Akshatha Laxmi (PES1201800130)

Sparsh Temani (PES1201800284)

Introduction/Background

Spin Locks: Reloops till the CAS operations returns True

- Equivalent of acquire lock: **while(!lock.CAS(0, 1));**
- Equivalent of lock release: **lock = 0**
- While the thread is in the critical section, lock is set to 1, once it exists, one of the waiting threads sets the lock to 1 again using the CAS operation and enters the critical section.
- CAS (Compare and Swap) : Executed as a single instruction on the CPU (atomic)
- Can have severe performance implications as only thread can enter the critical section at once.

Issues with Spin Lock based Concurrent Data Structures

Reloops on Progress and Non Progress

- Uses additional CPU cycles even when no progress is made
 - When a thread with an acquired lock gets preempted, another thread which gets scheduled will waste CPU cycles
- If the thread with an acquired lock dies, there is no progress made
 - If a thread holding a lock dies in the middle, other threads waiting for the lock cannot proceed
- Priority inversion
 - A process with a lower priority is holding a lock that is required by a process of a higher priority

How Lockless Programming is useful ?

- Lockless programming makes sure at least one thread is making progress at any given instant (assuming it is scheduled by the OS).
- When the thread is preempted, another process can make progress since no lock is acquired
- Since no lock is held in any of the threads, if a thread dies, only the operations assigned to the thread do not execute where as the rest of the system can proceed as it is
- Lockless Programming guarantees system wide progress

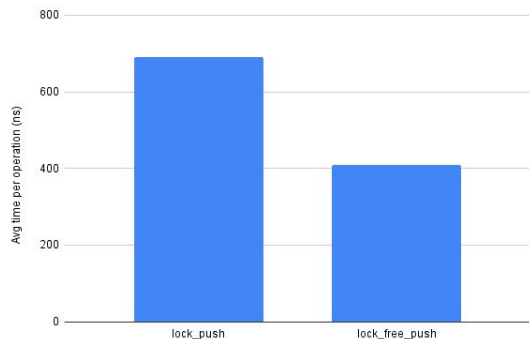
Implementation

- Language used: Go
- Lock and equivalent lock-free implementations benchmarked using Go testing package, and checked for race conditions using Go's race detector.
- Call Graphs generated using `pprof` for the lock-free and its respective lock-based counterparts
- Data Structures implemented for mid-term review:
 - **Stack**: Push, Pop, Peek
 - **Queue**: Enqueue, Dequeue
 - **List**: Insert, Delete


Quick Recap of Mid-term Review

Averaged over 100,000 runs


Stack (Push)




Queue (Enqueue)




List (Insert)




Stack (Pop)



Queue (Dequeue)



List (Delete)



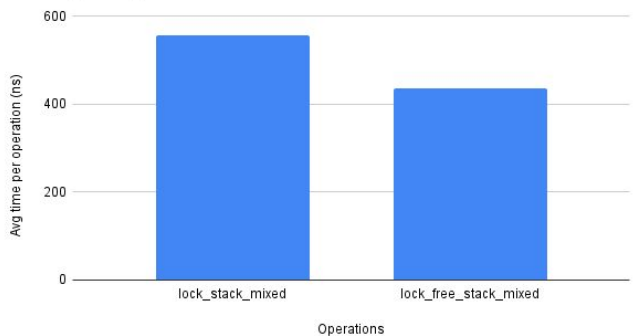
Progress Since Mid-term Review

- Improve implementation of delete operation in list
- Profile to see if false sharing is happening
 - Try and reduce false sharing via padding memory
 - Analyse its impact on performance
- Further implementation:
 - Lock free map
- Benchmark for read + write happening simultaneously


Benchmarks for concurrent Reads and Writes

Led to some (really) interesting realisations (end of ppt).


Stack (Mixed)




List (Mixed)



Queue (Mixed)



Cachelines and False Sharing



Profile for false sharing (and then cache miss)

- Use `perf c2c` to detect false sharing.

Shared Data Cache Line Table			(28 entries, sorted on Total HtMs)																	
Index	CacheLine	Node	PA cnt	Tot Hitm	Total	Load LclHitm	----- RmtHitm	Total records	Total Loads	Total Stores	----- L1Hit	----- L1Miss	----- FB	Core Load	Hit L1	----- L2	- LLC LclHit	Load Hit	-- LclHitm	-- RMT
0	0xc000016100	0	1	12.50%	4	4	0	7	5	2	2	0	0	0	0	0	1	4		
1	0xffff982f85628940	0	1	6.25%	2	2	0	3	3	0	0	0	0	0	0	0	1	2		
2	0xffff982f89e1d100	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
3	0xffff982f8aaf2880	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
4	0xffff982f95a0df80	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
5	0xffff982fcd75e0c0	0	4	3.12%	1	1	0	4	4	0	0	0	3	0	0	0	0	1		
6	0xffff9830b0b3fc80	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
7	0xffff9830e9e334c0	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
8	0xffff9830e9e6c580	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
9	0xffff9830e9e6cf00	0	1	3.12%	1	1	0	2	1	1	1	0	0	0	0	0	0	1		
10	0xffff9830e9eac580	0	4	3.12%	1	1	0	4	3	1	1	0	2	0	0	0	0	1		
11	0xffff9830e9eacd40	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
12	0xffff9830e9eec580	0	2	3.12%	1	1	0	2	1	1	1	0	0	0	0	0	0	1		
13	0xffff9830e9fb3500	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
14	0xffff9830e9fec500	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	1		
15	0xfffffddc333fce3ac0	N/A	0	3.12%	1	1	0	2	2	0	0	0	1	0	0	0	0	1		
16	0xc0000339800	0	2	3.12%	1	1	0	2	2	0	0	0	0	0	1	0	0	0	1	
17	0xc00002961c0	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
18	0xc00004fb800	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
19	0x7f373cc60c40	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
20	0x7f373dc9b740	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
21	0x7f373dcb000	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
22	0xc000024000	0	2	3.12%	1	1	0	2	2	0	0	0	0	0	0	0	1	1		
23	0xc000120a00	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
24	0x7fb026263f80	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
25	0x6235c0	0	3	3.12%	1	1	0	3	3	0	0	0	0	0	2	0	0	0	1	
26	0xc000123a40	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	
27	0xc000302700	0	1	3.12%	1	1	0	1	1	0	0	0	0	0	0	0	0	0	1	

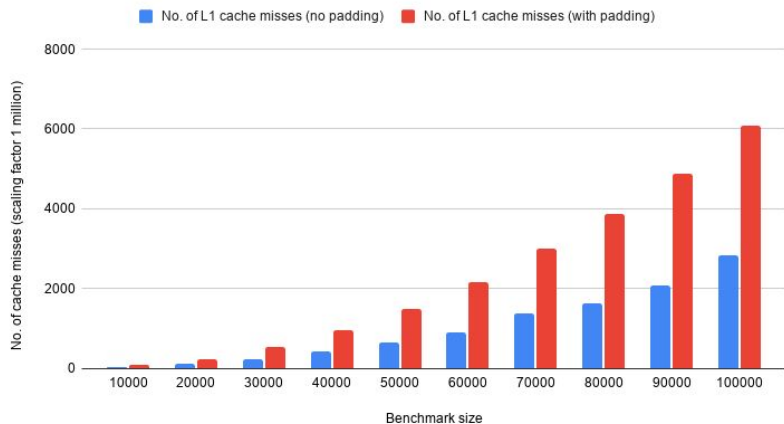
Profile for false sharing (and then cache miss)

Cacheline 0xc000016100																
----- HITM -----			-- Store Refs --			----- CL -----			----- cycles -----			Total	cpu			
RmtHitm	LclHitm	L1 Hit	L1 Miss	Off	Node	PA	cnt	Code address	rmt hitm	lcl hitm	load	records	cnt	Symbol	list.test	Shared Object
0.00%	100.00%	0.00%	0.00%	0x20	0		1	0x476c06	0	120	0	4	3	[.] 0x00000000000076c06	list.test	list.
0.00%	0.00%	100.00%	0.00%	0x20	0		1	0x476c2a	0	0	212	3	2	[.] 0x00000000000076c2a	list.test	list.

Profile for false sharing (and then cache miss)

- Use padding to try and avoid false sharing.
 - Side effects of doing so in the case of operations which are *read heavy*.
 - Validating results by profiling for cache misses using `perf stat`
 - `cat /sys/devices/system/cpu/cpu0/cache/index0/coherency_line_size`

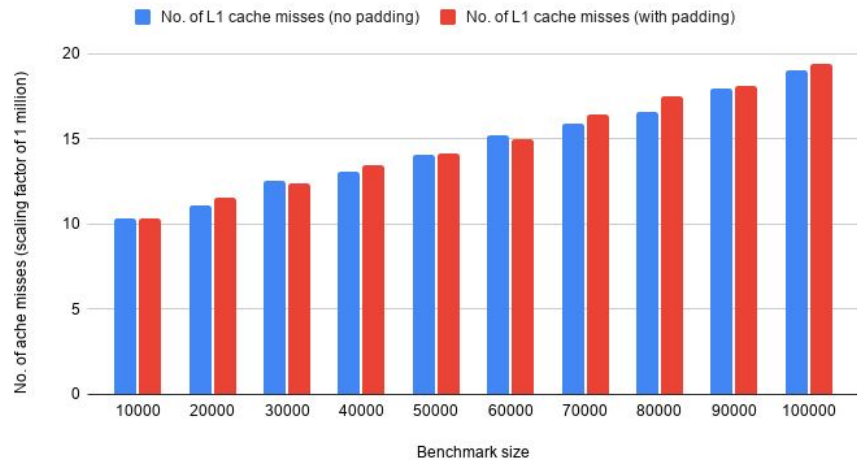
Cache miss comparison (List)



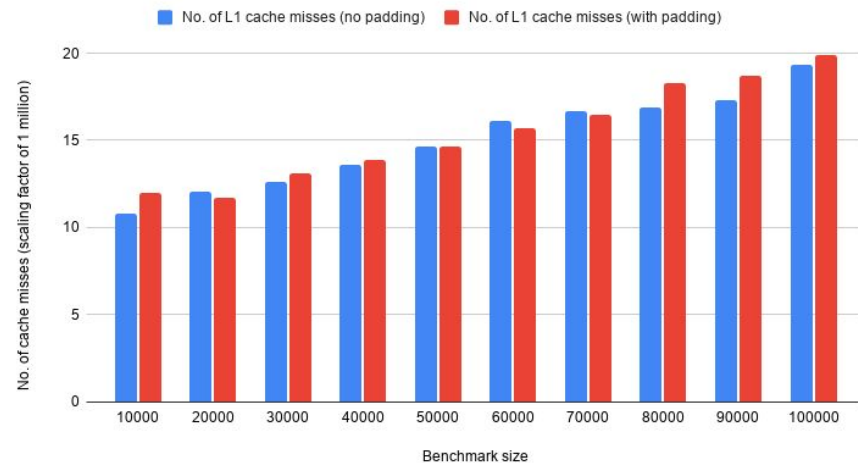
```
maddy@winston index0  
$ cat /sys/devices/system/cpu/cpu0/cache/index0/coherency_line_size  
64
```

Profiling for cache miss

Cache miss comparison (Stack)

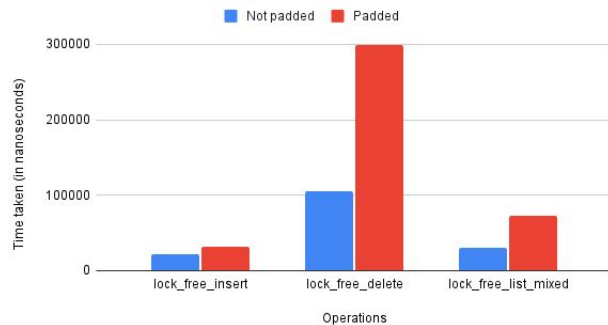


Cache miss comparison (Queue)

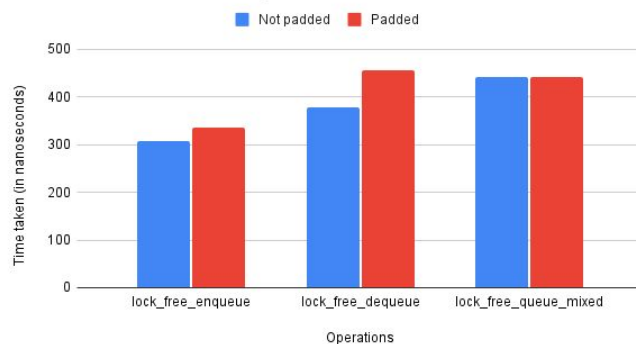


Operations with and without padding

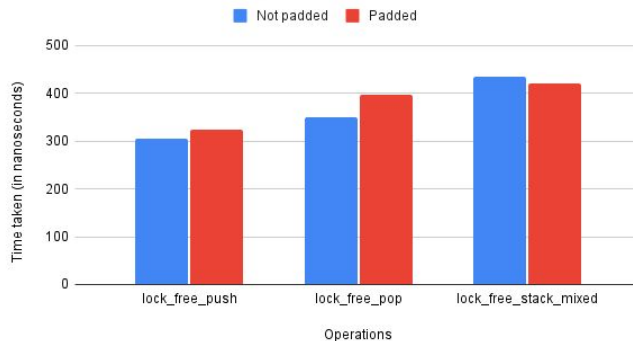
Time taken for different operations in Lock free List



Time taken for different operations in Lock free Queue



Time taken for different operations in Lock free Stack

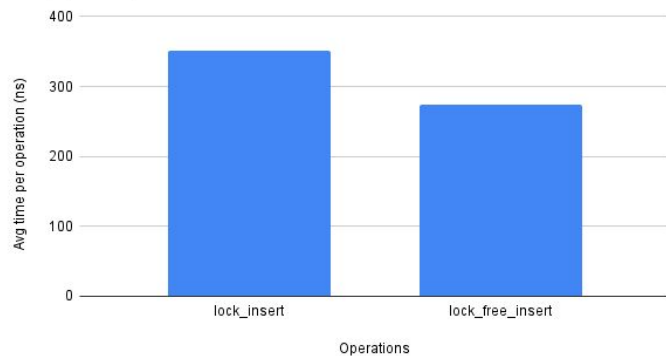


Lock Free Map (Methods)


- `Insert`: Insert the key specified in the map
- `InsertIfDoesntExist`: Inserts if the existing value is 'nil'
- `InsertCompare`: Takes a user defined function to decide if value should be inserted or not if the value already exists for the given key
- `Lookup`: If the element corresponding to a key exists, returns the element
- `Exists`: Checks if a key exists in the map or not

Lock free map

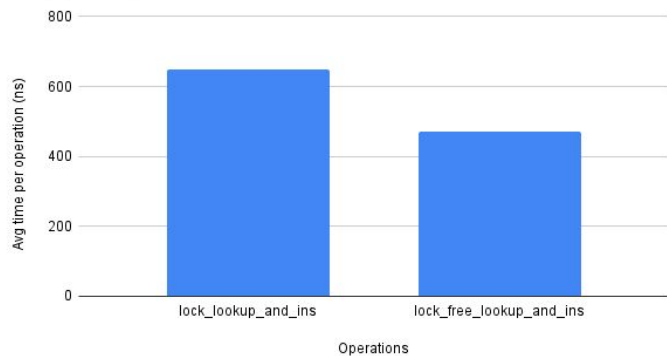
Map (Insert)



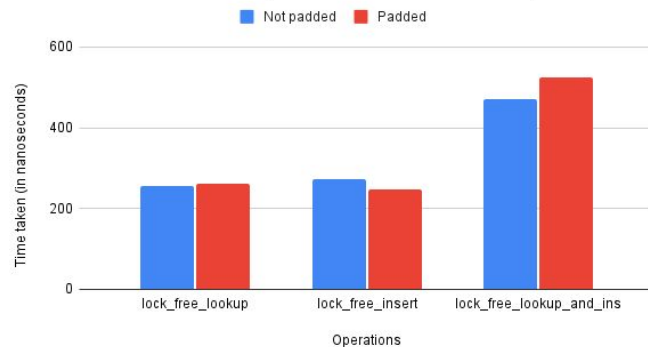
Map (Lookup)



Map (Mixed)

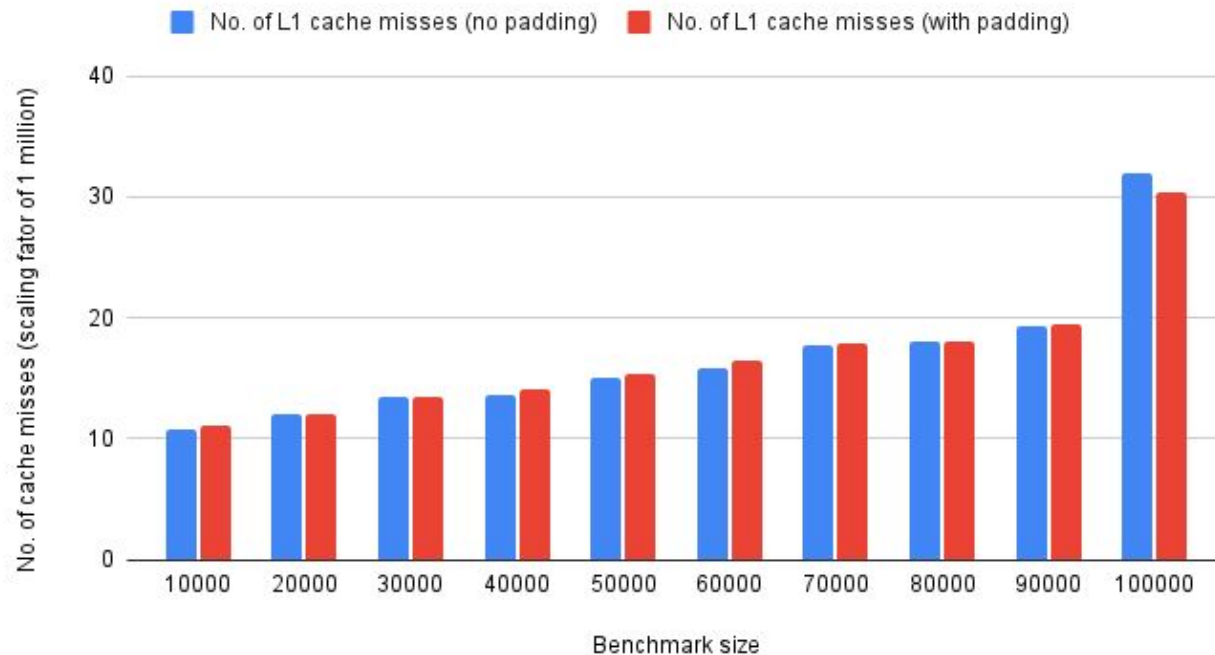


Time taken for different operations in Lock Free Map



Lock free map

Cache miss comparison (Map)



Miscellaneous Learnings

- Go race detector
 - Use of vector clocks for race condition detection
- Escape analysis
 - Can be seen using `-gcflags="-m"`
 - `"-m -m"` for more verbosity and so on.

```
./ops_test.go:74:19: int64(i) escapes to heap
./ops_test.go:82:19: int64(i) escapes to heap
./ops_test.go:95:19: BenchmarkLockDelAndIns ignoring self-assignment in queue.Tail = queue.Head
./ops_test.go:93:29: b does not escape
```

Work breakdown

Team member	Worked on	Time spent (approx.)
Sparsh Temani	Stack, Map	15 hours
Madhav Jivrajani	Queue, False sharing	15 hours
M S Akshatha Laxmi	List, profiling	15 hours

- Note: the above breakdown mostly signifies obtaining metrics/raw data and implementation of some form; deriving insights from these raw metrics and reasoning about performance based on implementation, was collectively done by the team.

References

- **Implementing Lock Free Queues**, *J. D. Valois, Dept. of CSE, Rensselaer Polytechnic Institute.*
- **A Pragmatic Implementation of Non-Blocking Linked-Lists**, *Timothy L. Harris, University of Cambridge.*
- **Introduction to Lock-free Programming** - Tony van Eerd, NDC Techtown
- **Lock Free Programming** - Herb Slutter, CppCon 2014
- **Designing a Lock-Free, Wait-Free Hash Map**, Shlomi Steinberg
- **A lock-free thread-safe HashMap optimized for fastest read access**, Cornel K
- **An intro to the Go race detector:** https://www.youtube.com/watch?v=4r9Kr_HtGdI
- **perf c2c:** <https://joemario.github.io/blog/2016/09/01/c2c-blog/>

Thank you